



# **Topic :** Valgrind - tool for memory debugging, memory leak detection, and profiling

Prepared By :

1. Sayantan Pandit (Exam Roll-111205022)

2. Suman Banerjee (Exam Roll-111205040)

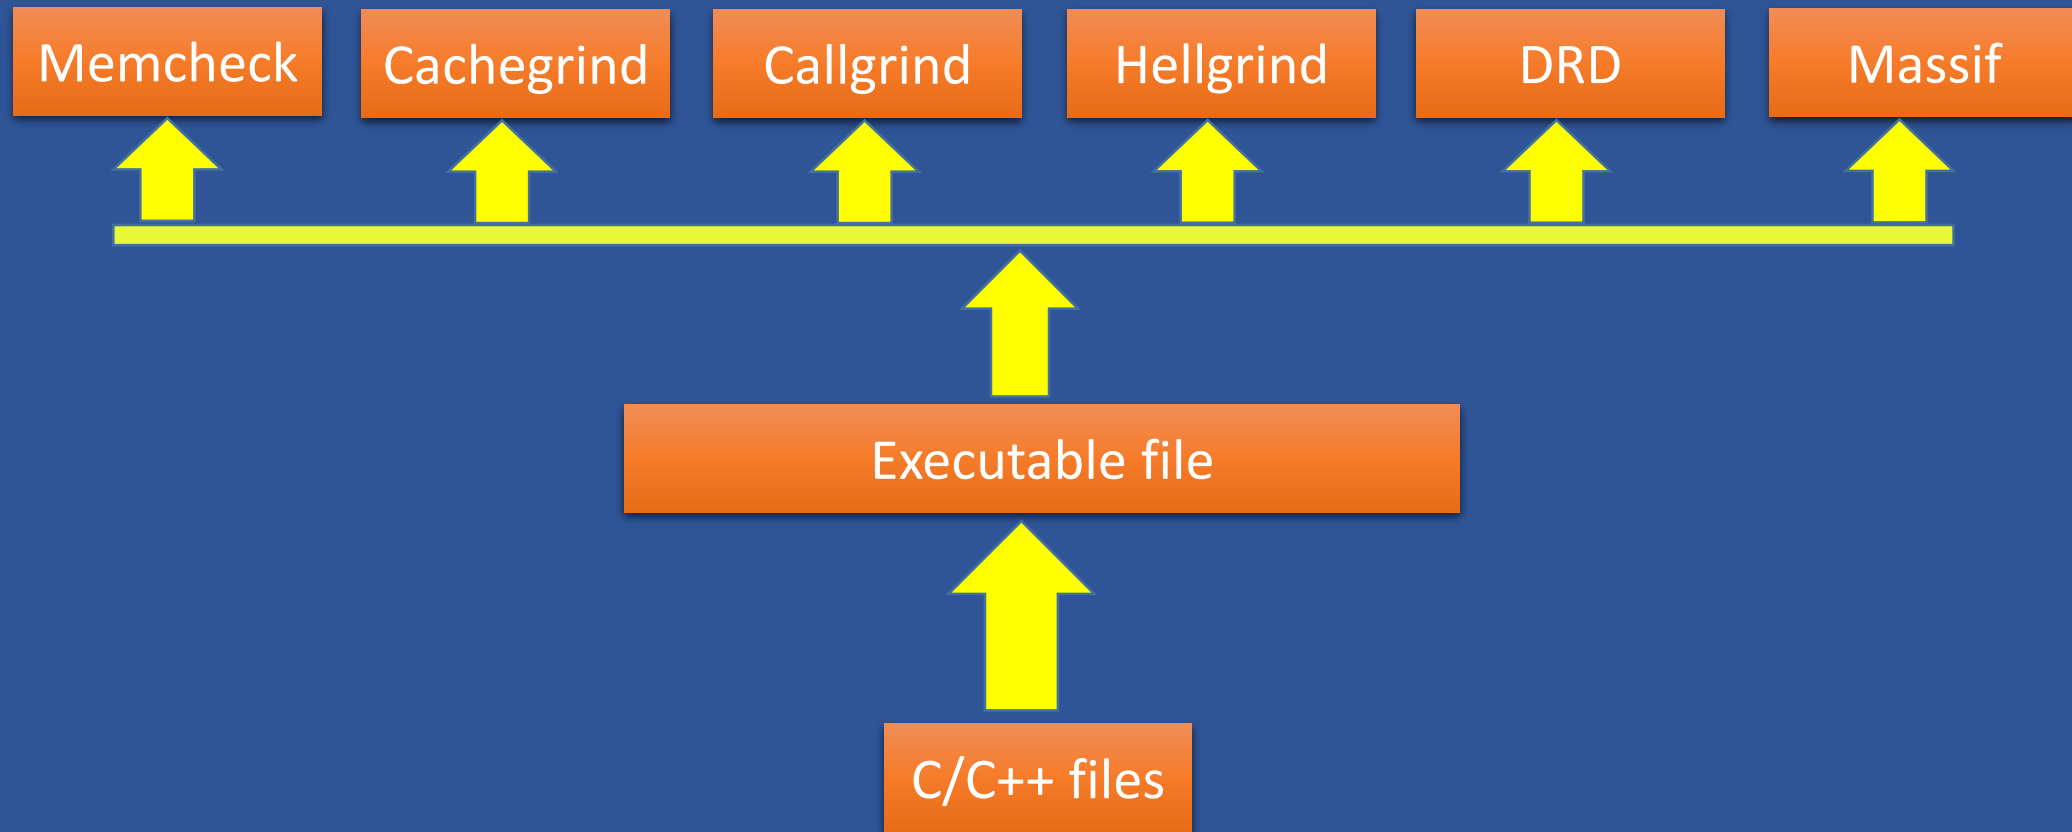
# INTRODUCTION

## A. What is Valgrind ?

1. Open-Source tool
2. Detects memory leaks/corruption
3. Provides a number of debugging and profiling tools.

# INTRODUCTION

## B. How does Valgrind work ?



# Installation of Software

## A. Directly from Repository

1. Open the terminal window
2. Type the command 'sudo apt-get install valgrind'
3. Enter your root password when prompted

## B. Download the Software

1. tar xvfz valgrind-1.0.0.tar.gz
2. cd valgrind-1.0.0
3. ./configure
4. make
5. make install

# Different Valgrind tools

1. Memcheck
2. Cachegrind
3. Callgrind
4. Helgrind
5. DRD
6. Massif
7. DHAT

# Memcheck

## Memory detected error

1. Use of Uninitialized values
2. Memory leaks detection
3. Invalid Pointer use
4. Accessing unavailable memory

# Memcheck

## A. Use of uninitialized values

```
int boo(int y)
{
    if(y == 2)
        printf("Correct\n");
}
int main()
{
    int x;
    boo(x);
}
```

```
rahul@rahul:~/college/sem7/lab/software/valgrind$ valgrind --tool=memcheck ./UnIntVariable
==9017== Memcheck, a memory error detector
==9017== Copyright (C) 2002-2013, and GNU GPL'd, by Julian Seward et al.
==9017== Using Valgrind-3.10.0.SVN and LibVEX; rerun with -h for copyright info
==9017== Command: ./UnIntVariable
==9017==
==9017== Conditional jump or move depends on uninitialised value(s)
==9017==   at 0x40053C: boo (UnIntVariable.c:5)
==9017==   by 0x40055B: main (UnIntVariable.c:12)
==9017==
==9017== Syscall param exit_group(status) contains uninitialised byte(s)
==9017==   at 0x4EF8309: _Exit (_exit.c:32)
==9017==   by 0x4E7321A: __run_exit_handlers (exit.c:97)
==9017==   by 0x4E732A4: exit (exit.c:104)
==9017==   by 0x4E58ECB: (below main) (libc-start.c:321)
==9017==
==9017== HEAP SUMMARY:
==9017==   in use at exit: 0 bytes in 0 blocks
==9017==   total heap usage: 0 allocs, 0 frees, 0 bytes allocated
==9017==
==9017== All heap blocks were freed -- no leaks are possible
==9017==
==9017== For counts of detected and suppressed errors, rerun with: -v
==9017== Use --track-origins=yes to see where uninitialised values come from
==9017== ERROR SUMMARY: 2 errors from 2 contexts (suppressed: 0 from 0)
```

# Memcheck

## B. Invalid Pointer Use

```
int main()
{
    char *x;
    x = (char *)malloc(10*sizeof(char));
    x[10] = 'a';
    free(x);
    return 0;
}
```

```
rahul@rahul:~/college/sem7/lab/software/valgrind$ valgrind --tool=memcheck ./InvPtrTest
==7941== Memcheck, a memory error detector
==7941== Copyright (C) 2002-2013, and GNU GPL'd, by Julian Seward et al.
==7941== Using Valgrind-3.10.0.SVN and LibVEX; rerun with -h for copyright info
==7941== Command: ./InvPtrTest
==7941==
==7941== Invalid write of size 1
==7941==    at 0x40059B: main (InvPtrTest.c:11)
==7941==   Address 0x51fc04a is 0 bytes after a block of size 10 alloc'd
==7941==    at 0x4C2AB80: malloc (in /usr/lib/valgrind/vgpreload_memcheck-amd64-linux.so)
==7941==   by 0x40058E: main (InvPtrTest.c:10)
==7941==
==7941==
==7941== HEAP SUMMARY:
==7941==    in use at exit: 0 bytes in 0 blocks
==7941==   total heap usage: 1 allocs, 1 frees, 10 bytes allocated
==7941==
==7941== All heap blocks were freed -- no leaks are possible
==7941==
==7941== For counts of detected and suppressed errors, rerun with: -v
==7941== ERROR SUMMARY: 1 errors from 1 contexts (suppressed: 0 from 0)
```



# Memcheck

## C. Accessing unavailable memory

```
int main()
{
    char *buf;
    buf = malloc(1<<32);
    fgets(buf, 1024, stdin);
    printf("s\n", buf);
    return 0;
}
```

```
==8051== by 0x4EA52C5: fgets (iofgets.c:56)
==8051== by 0x400631: main (SegTest1.c:11)
==8051== Address 0x51fc040 is 0 bytes after a block of size 0 alloc'd
==8051== at 0x4C2AB80: malloc (in /usr/lib/valgrind/vgpreload_memcheck-amd64-linux.so)
==8051== by 0x400615: main (SegTest1.c:9)
==8051==
==8051== Invalid write of size 1
==8051== at 0x4C2FD48: __GI_memcpy (in /usr/lib/valgrind/vgpreload_memcheck-amd64-linux.so)
==8051== by 0x4EA63B3: _IO_getline_info (iogetline.c:105)
==8051== by 0x4EA52C5: fgets (iofgets.c:56)
==8051== by 0x400631: main (SegTest1.c:11)
==8051== Address 0x51fc041 is 1 bytes after a block of size 0 alloc'd
==8051== at 0x4C2AB80: malloc (in /usr/lib/valgrind/vgpreload_memcheck-amd64-linux.so)
==8051== by 0x400615: main (SegTest1.c:9)
==8051==
==8051== Invalid write of size 1
==8051== at 0x4EA533A: fgets (iofgets.c:64)
==8051== by 0x400631: main (SegTest1.c:11)
==8051== Address 0x51fc042 is 2 bytes after a block of size 0 alloc'd
==8051== at 0x4C2AB80: malloc (in /usr/lib/valgrind/vgpreload_memcheck-amd64-linux.so)
==8051== by 0x400615: main (SegTest1.c:9)
```

# Cachegrind

Cachegrind is a tool for doing cache simulations and annotating the source line-by-line with the number of cache misses and interacts with a machine's cache hierarchy and (optionally) branch predictor.

# Cachegrind

```
#include <stdio.h>
#define N 1000
double array_sum(double a[][N]);
int main(int argc, char **argv) {
    double a[N][N];
    int i,j;
    for(i=0; i<N; i++) {
        for(j=0; j<N; j++)
            a[i][j] = 0.01;
    }
    printf("Sum = %10.3f\n", array_sum(a));
    return 0;
}

double array_sum(double a[][N]) {
    int i,j;
    double s;
    s=0;
    for(i=0; i<N; i++)
        for(j=0; j<N; j++)
            s += a[i][j];
    return s;
}
```

```
bmsd@ubuntu:~/Documents/7th_Sem/2nd_Valgrind$ valgrind --tool=cachegrind ./CacheGrind
==3004== Cachegrind, a cache and branch-prediction profiler
==3004== Copyright (C) 2002-2013, and GNU GPL'd, by Nicholas Nethercote et al.
==3004== Using Valgrind-3.10.0.SVN and LibVEX; rerun with -h for copyright info
==3004== Command: ./CacheGrind
==3004==
--3004-- warning: L3 cache found, using its data for the LL simulation.
Sum = 10000.000
==3004==
==3004== I   refs:      25,120,427
==3004== I1 misses:      926
==3004== LLi misses:      921
==3004== I1 miss rate:    0.00%
==3004== LLi miss rate:  0.00%
==3004==
==3004== D   refs:      13,048,524 (11,032,989 rd + 2,015,535 wr)
==3004== D1 misses:      251,884 ( 126,362 rd + 125,522 wr)
==3004== LLd misses:      251,709 ( 126,213 rd + 125,496 wr)
==3004== D1 miss rate:    1.9% ( 1.1% + 6.2% )
==3004== LLd miss rate:  1.9% ( 1.1% + 6.2% )
==3004==
==3004== LL refs:      252,810 ( 127,288 rd + 125,522 wr)
==3004== LL misses:      252,630 ( 127,134 rd + 125,496 wr)
==3004== LL miss rate:    0.6% ( 0.3% + 6.2% )
```

# Massif

Massif is a heap profiler. It measures how much heap memory the program uses. This includes both the useful space, and the extra bytes allocated for book-keeping and alignment purposes. It can also measure the size of the program's stack(s).

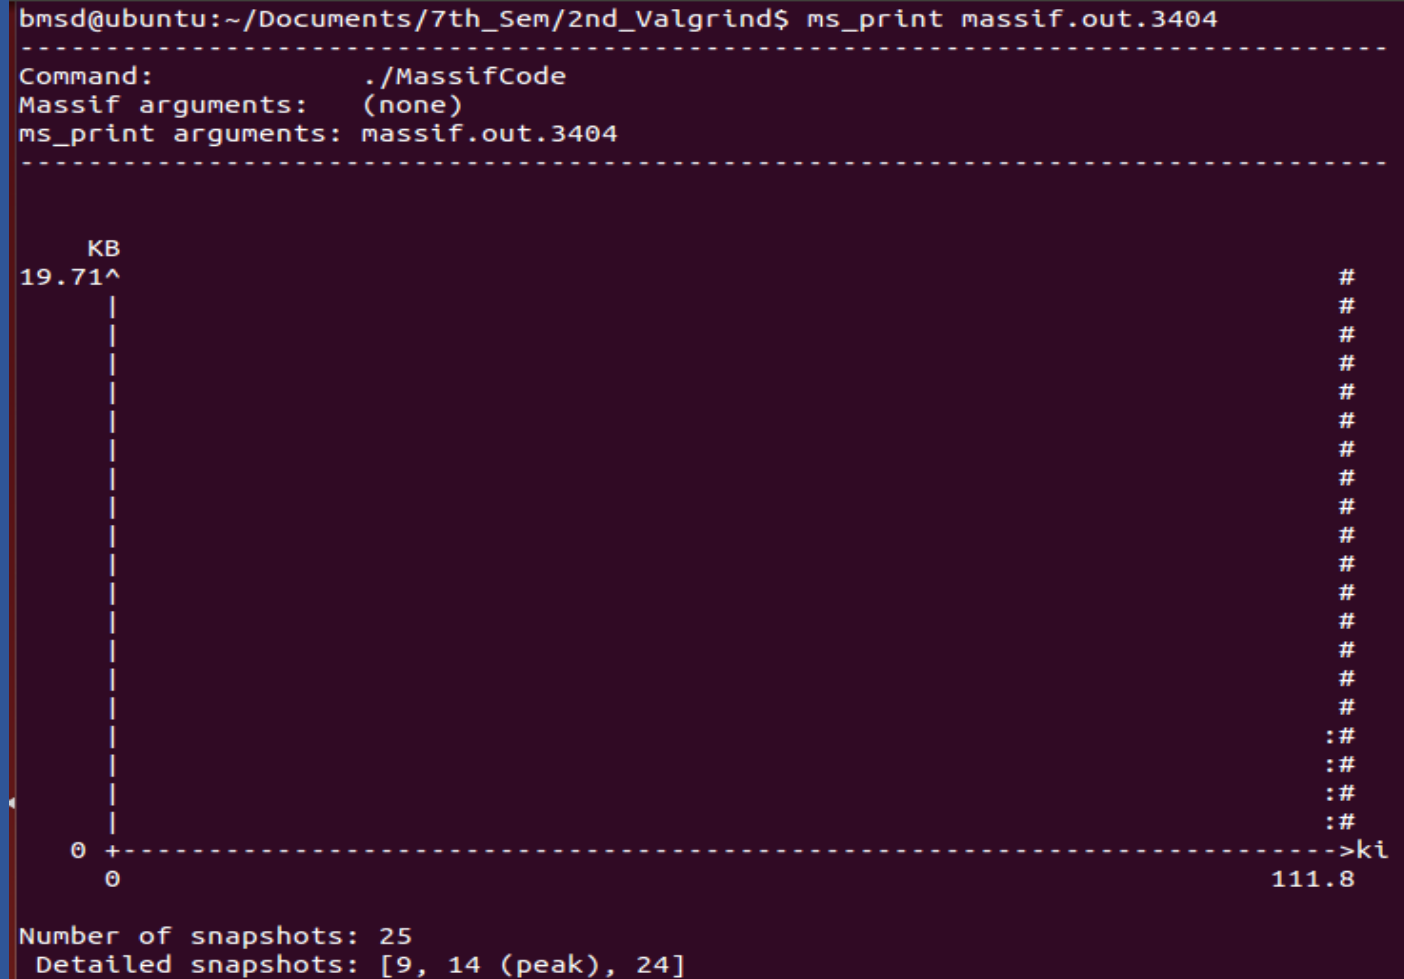
# Massif

```
#include <stdio.h>

void g(void) {
    malloc(4000);
}

void f(void) {
    malloc(2000);
    g();
}

int main(void) {
    int i;
    int *a[10];
    for(i=0; i<10; i++)
        a[i] = malloc(1000);
    f();
    g();
    for(i=0; i<10; i++)
        free(a[i]);
    return 0;
}
```



# Massif

n	time(i)	total(B)	useful-heap(B)	extra-heap(B)	stacks(B)
0	0	0	0	0	0
1	112,814	1,016	1,000	16	0
2	112,855	2,032	2,000	32	0
3	112,896	3,048	3,000	48	0
4	112,937	4,064	4,000	64	0
5	112,978	5,080	5,000	80	0
6	113,019	6,096	6,000	96	0
7	113,060	7,112	7,000	112	0
8	113,101	8,128	8,000	128	0
9	113,142	9,144	9,000	144	0
98.43% (9,000B) (heap allocation functions) malloc/new/new[], --alloc-fns, etc. ->98.43% (9,000B) 0x4005BB: main (MassifCode.c:18)					
n	time(i)	total(B)	useful-heap(B)	extra-heap(B)	stacks(B)
10	113,183	10,160	10,000	160	0
11	113,227	12,168	12,000	168	0
12	113,264	16,176	16,000	176	0
13	113,305	20,184	20,000	184	0
14	114,199	20,184	20,000	184	0
99.09% (20,000B) (heap allocation functions) malloc/new/new[], --alloc-fns, etc. ->49.54% (10,000B) 0x4005BB: main (MassifCode.c:18)   ->39.64% (8,000B) 0x400589: g (MassifCode.c:4)   ->19.82% (4,000B) 0x40059E: f (MassifCode.c:10)     ->19.82% (4,000B) 0x4005D7: main (MassifCode.c:20)       ->19.82% (4,000B) 0x4005DC: main (MassifCode.c:21)   ->09.91% (2,000B) 0x400599: f (MassifCode.c:9) ->09.91% (2,000B) 0x4005D7: main (MassifCode.c:20)					

# Helgrind

Helgrind is a thread error detector.

## A Simple Data Race :

```
#include <stdio.h>
#include <pthread.h>
int var = 0;
void* child_fn ( void* arg ) {
    var++;    /* Unprotected relative to parent */ /* this is line 6 */
    return NULL;
}
int main ( void ) {
    pthread_t child;
    pthread_create(&child, NULL, child_fn, NULL);
    var++; /* Unprotected relative to child */ /*this is line 13 */
    pthread_join(child, NULL);
    return 0;
}
```

# Limitation

1. Program runs 20 to 50 times slower
2. Increased memory consumption
3. Highly optimized code may cheat Valgrind



**THANK YOU**